

Perception check: the model perceives ~everything (Pass-1 $\approx$ 1.0) but end-to-end reasoning drops ~35%

